

# ECE 4301 Midterm Report: Encrypted Video Streaming

## Design

Our encrypted video streaming system comprises two major components: the GStreamer video pipeline, and the encrypt/decrypt module.

On the sender side, the data flows from a Video4Linux2 source through a decodebin element which turns the picture series from the webcam into a video format. It continues through an h264 encoder which compresses the video stream to reduce the work required by the encryption module. GStreamer channel capabilities are added using the GDP protocol. Then the data arrives at an Appsink element where it is picked up by the encrypter module. The encrypter module polls the Appsink element for a sample, and when it is ready the module encrypts the sample and sends the ciphertext over the network to the receiver side.

On the receiver side, the ciphertext is received from the sender over the network. The receiver decrypts the payload and then pushes the plaintext to an Appsrc element at which point the data re-enters the gstreamer pipeline. Channel capabilities are discovered by a gdpdepay element, and then the video is forwarded to an h264 decoder to reproduce the original video stream. This video stream goes into an fpsdisplaysink element which renders the video feed and the FPS of the video.

The encrypt/decrypt module additionally frames data into packets, allowing data packets to be freely intermixed with control packets to support live rekeying.

The program is configured at execution time with command line flags to set the program mode—“sender” or “receiver”—the asymmetric encryption mode (“rsa” for RSA or “ecdh” for Elliptic-Curve Diffie-Hellman) and the IP address and port of the connection.

It was written entirely in Rust, using the aes-gcm crate for steady-state streaming encrypt/decrypt, which can be compiled with or without hardware acceleration, allowing us to easily examine the effects of hardware acceleration.

## Protocol Diagram

The packet protocol for all packets sent by the program conformed to the protocol pictured in the following image. All packets start with a predetermined start byte, followed by a byte to specify the packet type. Currently, this protocol allows four types of packets. These types are (1) general, encrypted packets; (2) public key packets; (3) encrypted session key packets; and (4) a rekey request packet (to initiate the cycling out of an old symmetric encryption key). Following this, the length of the data portion of the packet is specified in bytes with an unsigned, 32-bit integer. The data portion and the system time portion follow this. The system time portion of the packet has a length of 17 bytes. Finally, all packets end with an end byte.

Start Byte (u8)	Packet Type (u8)	Data Length (u32)	Packet data (Vec<u8>)	System Time (Vec<u8>)	End Byte (u8)
--------------------	------------------	-------------------	--------------------------	--------------------------	------------------

## Security Choices

For the asymmetric encryption handshake, the program can use either RSA 3072 to send a session key or the Elliptic Curve Diffie-Hellman (ECDH) key-exchange protocol. In the case of the RSA 3072 handshake, the transmitter sends the session key to the receiver via the receiver's public key. On each rekey request, the transmitter generates a new session key, and the receiver generates a new public key. In the case of ECDH, both generate new public keys. In this application, old keys are not stored. Further, keys in use are wiped when the application exits.

Once a symmetric key is shared between the transmitter and receiver, data packets, encrypted with AES-GCM encryption, are sent between the transmitter and receiver. The packets that are sent are authenticated (for AEAD support) with a packet number that is equal to the number of packets that were sent before it. This number is used in the encryption of each packet. When the receiver receives a packet of data, it must use this number to successfully decrypt the packet.

## Measurement Setup

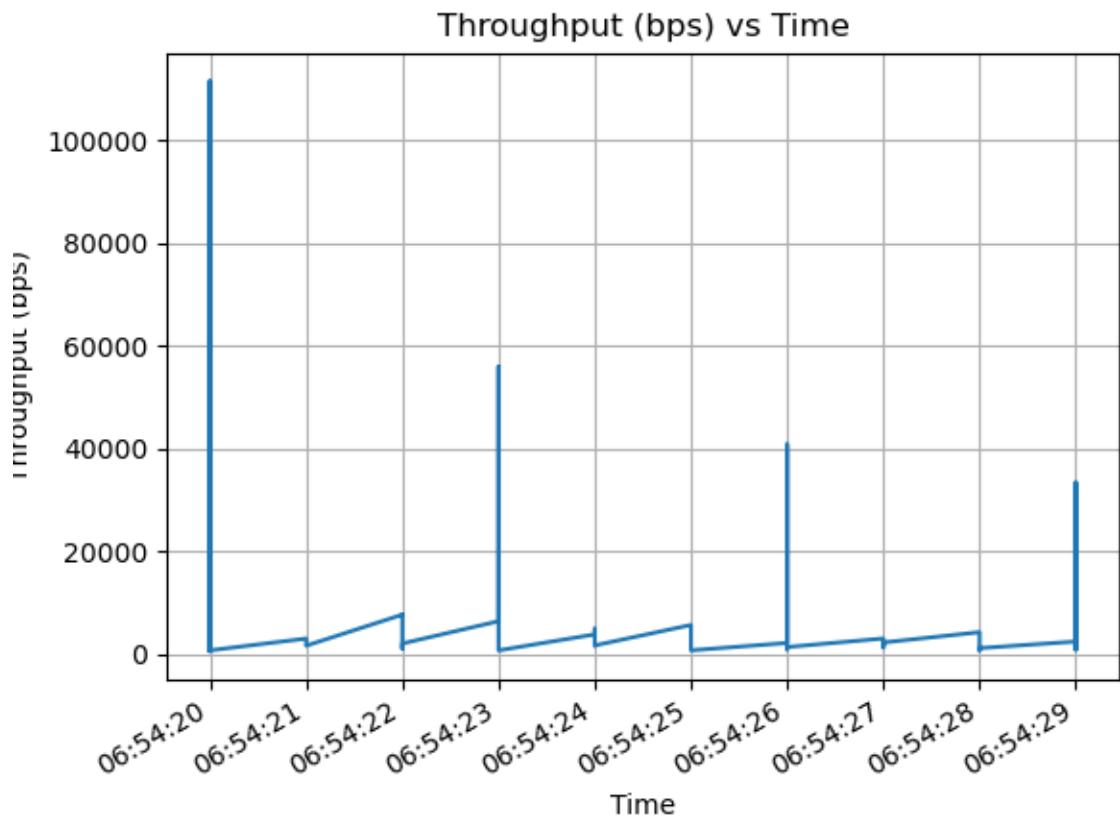
A Klein Tools ET920 inline power meter was used to acquire power measurements from the Raspberry Pi. An image of this setup follows.




## Results

We were able to achieve a framerate of 30fps with a 640p camera feed without any significant jitter. We observed that ECDH key exchange was very efficient, finishing in usually under a millisecond, while RSA key exchange could take up to a half a second. We did not observe a significant change in power usage depending on whether hardware acceleration was turned on or off. Following are graphs of our data.

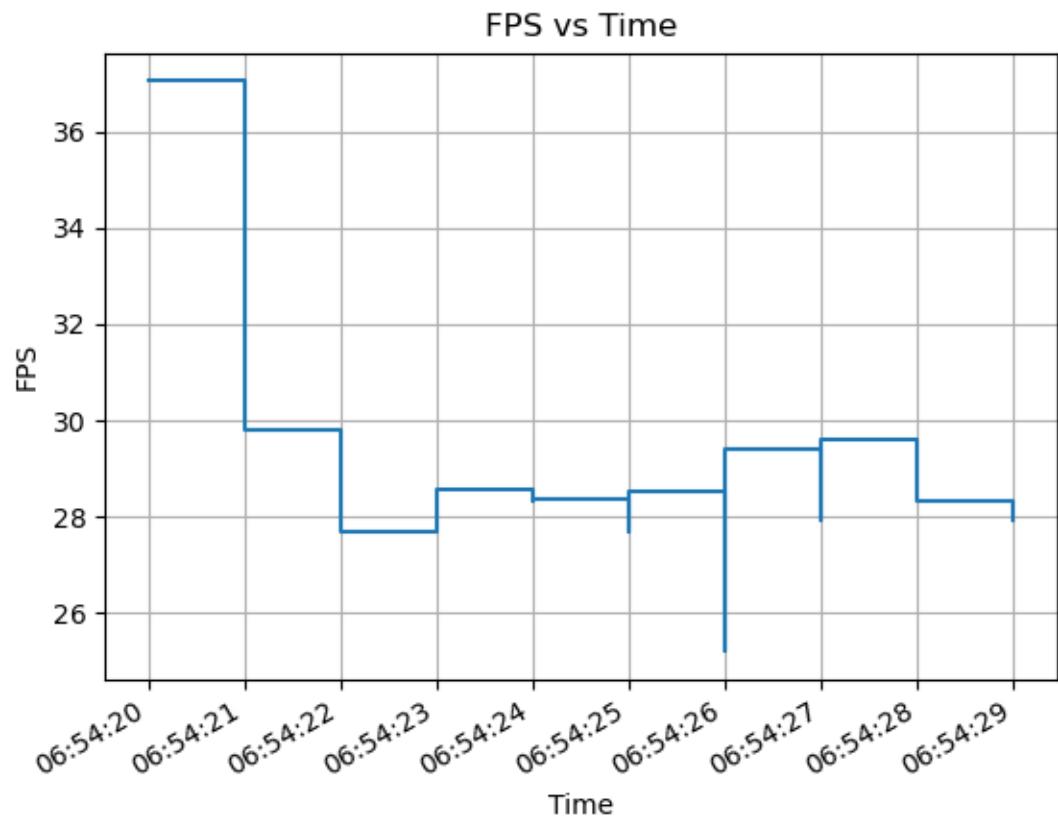
## Steady State Throughput



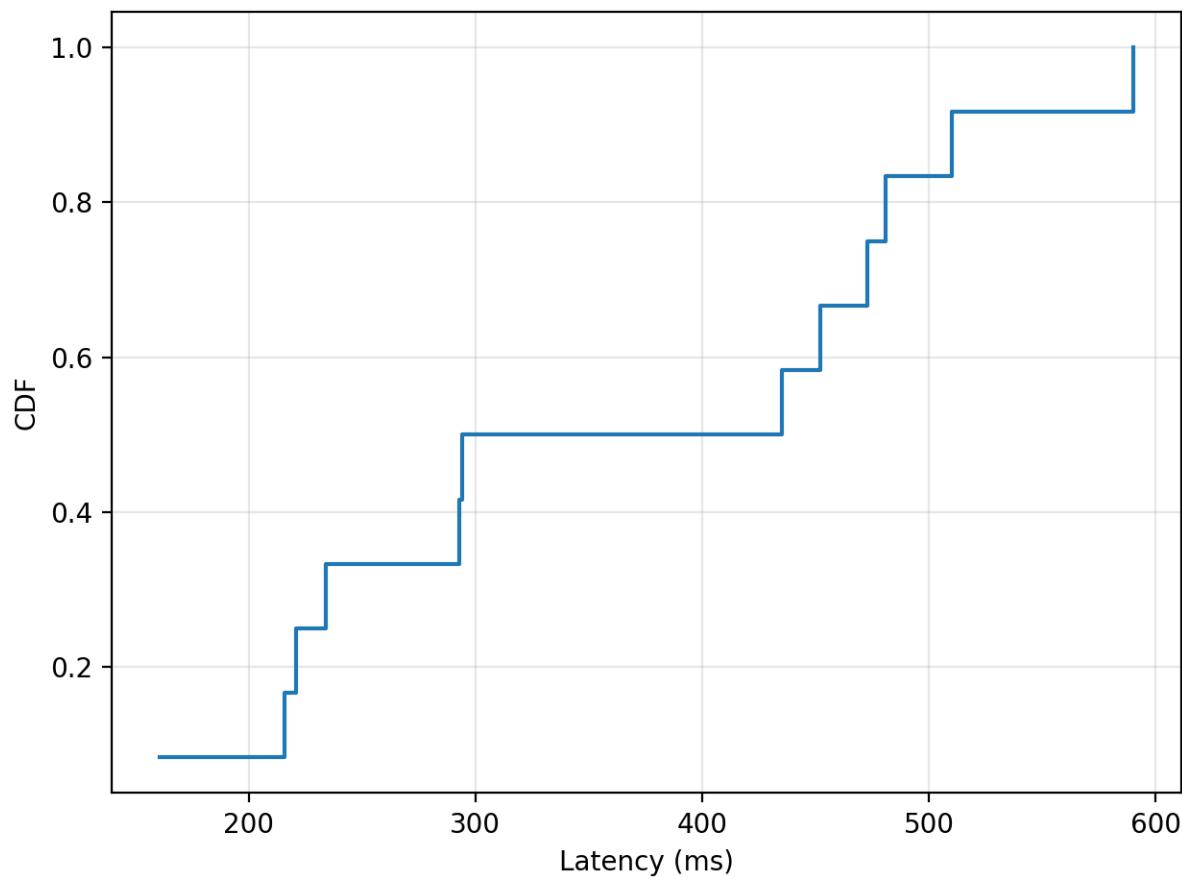
## ECDH Latency



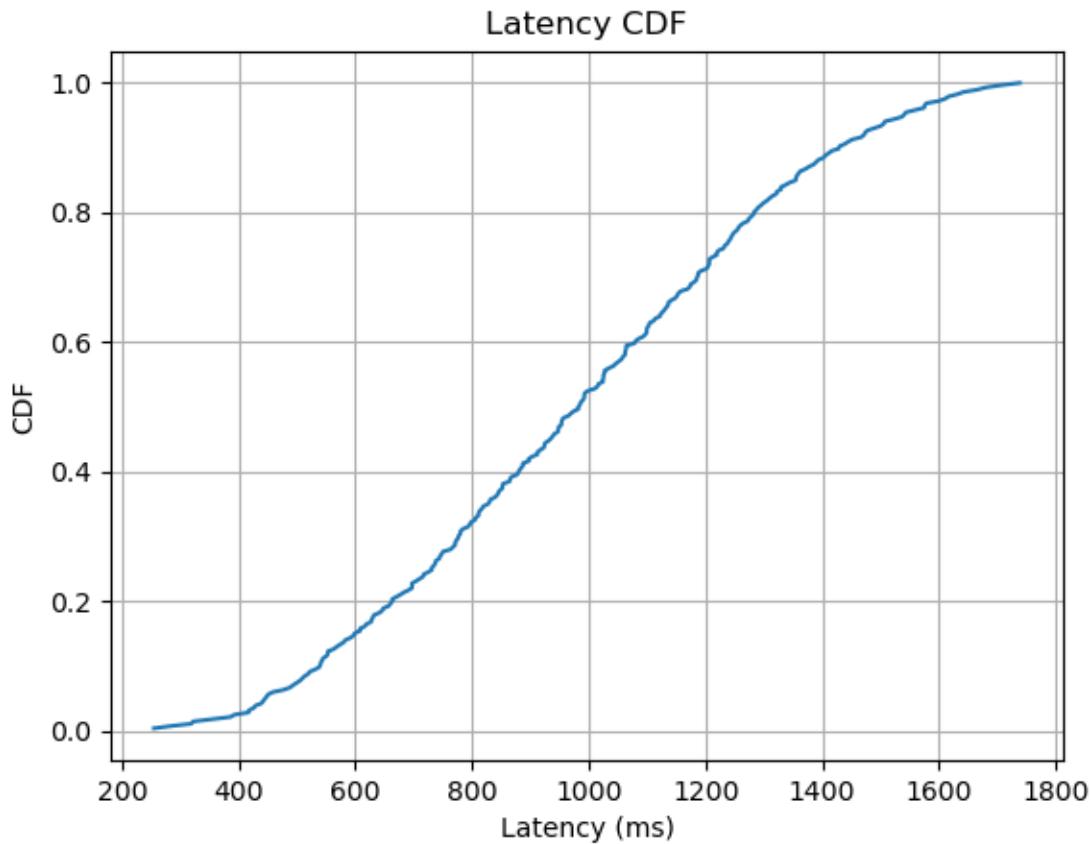
## Steady State FPS



## RSA Latency



## Steady State Latency



## Discussion

Because we did not observe any significant changes in throughput or execution time with respect to enabling or disabling hardware acceleration, we believe that the Pi's cryptographic extension does not provide significant value to this type of application. This is likely because, for this sort of application, the video processing computations dominate the Pi's CPU. So, the cryptographic hardware acceleration on the Raspberry Pi is likely of *limited* utility.